

# Activity 1

## Part 1

The screenshot shows the Apache Spark 4.0.0 UI interface for the 'Jobs' tab. At the top, there are navigation links for 'Jobs', 'Stages', 'Storage', 'Environment', 'Executors', 'SQL / DataFrame', and 'Structured Streaming'. On the right, there are icons for 'Anmelden', a download arrow, a '1' notification, a '2' notification, and a menu icon. The main content area is titled 'Spark Jobs' with a question mark icon. It displays the following information:

- User:** spark
- Started At:** 2026/01/17 15:30:13
- Total Uptime:** 2.4 min
- Scheduling Mode:** FIFO
- Active Jobs:** 1
- Completed Jobs:** 4

Below this, there are links for 'Event Timeline' and 'Active Jobs (1)'. The 'Active Jobs (1)' section shows a table with one row:

Job Id (Job Group) ▾	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
4 (7b540083-dc25-4f73-a160-c3295e36e752)	id = c7171bd0-6c26-47b0-9238-657df6b646cd runId = 7b540083-dc25-4f73-a160-c3295e36... start at <unknown>:0	2026/01/17 15:32:11 (kill)	24 s	1/2	189/202 (2)




At the bottom of the table, there are pagination controls: 'Page: 1', '1 Pages. Jump to 1', 'Show 100 items in a page.', and a 'Go' button.


## Part 2

a



## ▼ DAG Visualization






One with skip and one without it.

b

Shuffle Read	Shuffle Write
12.6 KiB	
	12.6 KiB
7.9 KiB	
12.6 KiB	7.9 KiB
12.6 KiB	
	12.6 KiB
7.9 KiB	
12.7 KiB	7.9 KiB

shuffle write found in a stage



Example of how the Tasks then look like if you hover over just one of them as I found it interesting.

```

Task 0 (attempt 0)
Status: SUCCESS
Launch Time: 2026/01/17 15:45:32
Scheduler Delay: 5 ms
Task Deserialization Time: 8 ms
Shuffle Read Time: 0 ms
Executor Computing Time: 2 s
Shuffle Write Time: 25 ms
Getting Result Time: 0 ms
Result Serialization Time: 0 ms
Getting Result Time: 0 ms

```

C

For right now we only have one worker and one driver, so the worker has all the work. But from what I understand this is how it should be for one worker and one driver.

	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
Active(2)	0	17.7 MiB / 848.3 MiB	0.0 B	1	2	0	11770	11772	33 min (11 s)	0.0 B	774.9 KiB	484.6 KiB	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(2)	0	17.7 MiB / 848.3 MiB	0.0 B	1	2	0	11770	11772	33 min (11 s)	0.0 B	774.9 KiB	484.6 KiB	0

**Executors**

Show 20 entries Search:																			
Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Logs	Thread Dump	Heap Histogram	Add Time	Remove Time
driver	b9ec450fed29:42325	Active	0	8.9 MiB / 434.4 MiB	0.0 B	0	0	0	0	0	17 min (1 s)	0.0 B	0.0 B	0.0 B		Thread Dump	Heap Histogram	2026-01-17 16:30:15	-
0	172.19.0.5:46657	Active	0	8.9 MiB / 413.9 MiB	0.0 B	1	2	0	11770	11772	16 min (10 s)	0.0 B	774.9 KiB	484.6 KiB	stdout stderr	Thread Dump	Heap Histogram	2026-01-17 16:30:22	-

## Part 3

1


Stage 168 and 113 for me with 35 seconds.

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
168	id = c7171bd0-6c26-47b0-9238-657df6b646cd runId = 7b540083-dc25-4f73-a160-c3295e36e752 batch = 33 start at <unknown>:0	2026/01/17 15:51:40	35 s	200/200			12.6 KiB	8.0 KiB
113	id = c7171bd0-6c26-47b0-9238-657df6b646cd runId = 7b540083-dc25-4f73-a160-c3295e36e752 batch = 22 start at <unknown>:0	2026/01/17 15:44:57	35 s	200/200			12.6 KiB	7.9 KiB

From what I can tell our biggest bottleneck right now is the limited parallelism.



These two pictures are event Timelines from 2 of the slowest ones.




This one was done with the same amount of tasks in less then half of the time.

But as you can see it has like one 2 long bars running at the same time and lots of short ones scattered in the beginning.

We only added 1 Executor Core, 1 Executor and 1 GB of memory.

2



So in total it is right now using 22.7 MiB out of 848.3 MiB which are available to it

3

In Spark Structured Streaming, performance and scalability are about how well the system processes streaming data and how it handles growing workloads.

Structured Streaming usually works with micro-batches, meaning incoming data is processed in small chunks. Smaller batches give lower latency but more overhead, while larger batches increase throughput but also delay results.

Performance depends on how efficiently Spark uses cluster resources like CPU and memory. Operations such as joins, aggregations, and stateful processing can slow things down if they require a lot of memory or shuffling.

Scalability means Spark can handle higher data rates by adding more executors or increasing parallelism. However, bottlenecks like network shuffles, state size, and slow sinks can limit how much it actually scales.

## Activity 2

---

change to the spark call

```
spark-submit \
  --master spark://spark-master:7077 \
  --packages org.apache.spark:spark-sql-kafka-0-10_2.13:4.0.0 \
  --num-executors 2 \
  --executor-cores 2 \
  --executor-memory 2G \
  /opt/spark-apps/spark_structured_streaming_logs_processing.py
```

Changes to docker file in Exercise 3

```
spark-worker:  
  image: bitnamilegacy/spark:4.0.0  
  depends_on:  
    - spark-master  
  environment:  
    - SPARK_MODE=worker  
    - SPARK_MASTER_URL=spark://spark-master:7077  
    - SPARK_WORKER_CORES=4  
    - SPARK_WORKER_MEMORY=4G  
  networks:  
    - streaming-net  
  deploy:  
    replicas: 1  
    resources:  
      limits:  
        memory: 4095M  
        cpus: '4'
```

Changes to Docker file in Load-generator

```
services:  
  generator:  
    image: adrianovogel/hgb-load-gen:latest  
    extra_hosts:  
      - "host.docker.internal:host-gateway"  
    environment:  
      - KAFKA_BROKER=host.docker.internal:9095  
      - KAFKA_TOPIC=logs  
      - TARGET_RPS=10000  
      - ADDITIONAL_TERM=crash  
      - ADDITIONAL_TERM_RATE=100  
    deploy:  
      replicas: 4 # This enables to run more instances (containers)  
      resources:  
        limits:  
          memory: 2048M  
          cpus: '1'  
      networks:  
        - streaming-net
```

## Monitoring

Input Rate vs. Process Rate

35671.60	61106.03

I think I do not have enough avg Input per second right now but it at least shows me that it works (I was a little bit scared that this all would be to much for my laptop, that is why it might be like that.)

## The Executor Tab

The lovely flame chart from the Threat dump




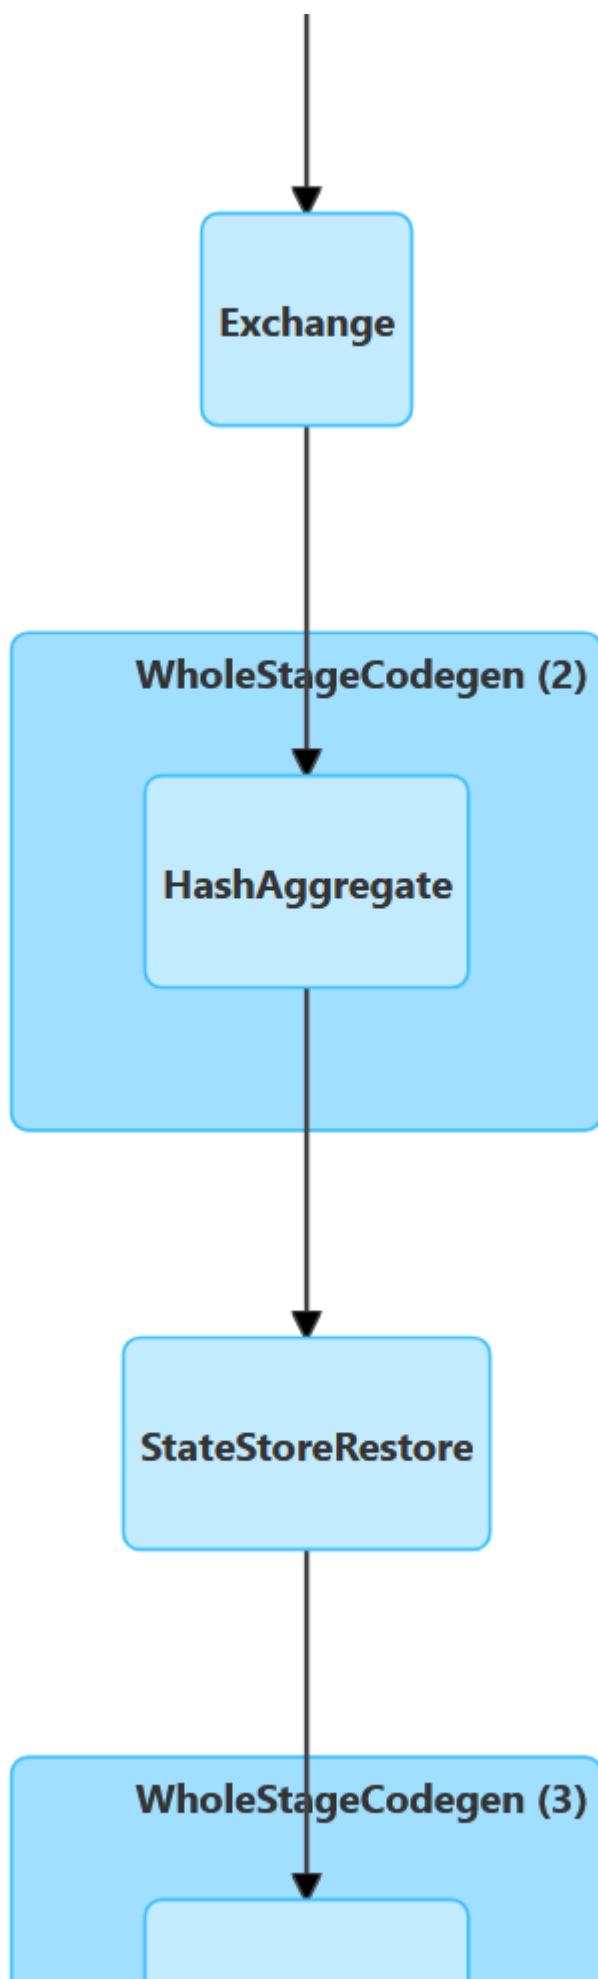
And the tasks


Complete Tasks	Total Tasks
0	0
5810	5810
5800	5800


I think they are quite evenly distributed which is good. Again the one at the top is the driver therefore this does not count.

The SQL/Query Tab









This is sadly the best way I am able to show the DAG as it is just a line downwards.

It has two exchange blocks. I cannot fully state if this is now being evenly distributed across all of my cores, but I can say if there is a problem it is in the exchange as I read in the details down:

(5) Exchange Input [2]: [source\_ip#18, count#32L] Arguments: hashpartitioning(source\_ip#18, 200), REQUIRED\_BY\_STATEFUL\_OPERATOR, [plan\_id=15493]

Which is also true for the other partitions. Now if one source IP is more frequent than the other (the one thing I cannot tell you right now as I do not know how I would figure that out), then it would be distributed to this again, so small number of partitions receive the most data would be the skewness I think. But right now we only have 2 Partitions and I do not know how I can check this.